

Getting Data from Multiple Tables – The JOIN

- In order to run a query that retrieves data from multiple tables, those tables can be **JOINED**.

Getting Data from Multiple Tables – The JOIN

- In order to run a query that retrieves data from multiple tables, those tables can be **JOINED**.
- Joining two tables requires that the two tables have a common key (typically a foreign key relationship) that appears in both tables.

Getting Data from Multiple Tables – The JOIN

- In order to run a query that retrieves data from multiple tables, those tables can be **JOINED**.
- Joining two tables requires that the two tables have a common key (typically a foreign key relationship) that appears in both tables.
- The common key columns need NOT have the same name, but must be of the same data type and length.

Getting Data from Multiple Tables – The JOIN

- In order to run a query that retrieves data from multiple tables, those tables can be **JOINED**.
- Joining two tables requires that the two tables have a common key (typically a foreign key relationship) that appears in both tables.
- The common key columns need NOT have the same name, but must be of the same data type and length.
- A JOIN is one of the most **resource intensive** activities one can do in a relational database.

Basic Example

Let's join nwOrders to nwEmployees

- nwOrders has 830 rows, each with an EmployeeID
- nwEmployees has 9 rows, each with an EmployeeID
- They have a common key: EmployeeID (primary key in nwEmployees; foreign key in nwOrders)
- We want SQL to join the rows in nwEmployees and nwOrders where the EmployeeID matches

SQL Joins

Provide a listing showing Northwinds employees, sorted by LastName, and a count of each employee's orders from highest to lowest

```
Select LastName, Firstname, count(OrderID) as 'Orders'  
from nwEmployees, nwOrders  
where nwEmployees.EmployeeID =  
      nwOrders.EmployeeID  
GROUP BY LastName, FirstName  
Order By 1, 3 desc
```

Qualifying the Column Names

- Since the column “EmployeeID” exists in BOTH tables in this query, when referring to EmployeeID, we need to tell SQL which one.
- Therefore, we suffix the table name in front of the column name separated by a “.”
- Failure to fully qualify the column name will result in an “ambiguous column” error

Alternative

To save some typing, we can define an “alias” for each table. We can temporarily – only for the duration of this query -- rename the nwEmployees table “E”, and rename the nwOrders table “O”.

```
Select LastName, Firstname, count(OrderID) as 'Orders'  
      from nwEmployees E, nwOrders O  
     where E.EmployeeID = O.EmployeeID  
GROUP BY LastName, FirstName  
Order By 1
```

Beware the Cartesian Product

- Product = one table multiplied by another table
- The JOIN often creates a product, then selects the product where the keys match
- For example, let's join nwOrders to nwEmployees
- nwOrders has 14 columns, 830 rows
- nwEmployees has 17 columns, 9 rows
- The Cartesian product has 31 (14+17) columns, and 7470 (830 * 9) rows – most of which are meaningless



Cartesian Product

Stopped here Friday 2/22

SQL Joins

Cartesian Product from an unqualified join:

customerid	companyname	contactname	country	OrderID	CustomerID	Orderdate	shipcountry
GREAL	Great Lakes Food Market	Howard Snyder	USA	10262	RATTC	2013-07-22	USA
HUNGC	Hungry Coyote Import Store	Yoshi Latimer	USA	10262	RATTC	2013-07-22	USA
LAZYK	Lazy K Kountry Store	John Steel	USA	10262	RATTC	2013-07-22	USA
LETSS	Lets Stop N Shop	Jaime Yorres	USA	10262	RATTC	2013-07-22	USA
LONEP	Lonesome Pine Restaurant	Fran Wilson	USA	10262	RATTC	2013-07-22	USA
OLDWO	Old World Delicatessen	Rene Phillips	USA	10262	RATTC	2013-07-22	USA
RATTC	Rattlesnake Canyon Grocery	Paula Wilson	USA	10262	RATTC	2013-07-22	USA
SAVEA	Save-a-lot Markets	Jose Pavarotti	USA	10262	RATTC	2013-07-22	USA
SPLIR	Split Rail Beer & Ale	Art Braunschweiger	USA	10262	RATTC	2013-07-22	USA
THEBI	The Big Cheese	Liz Nixon	USA	10262	RATTC	2013-07-22	USA
THECR	The Cracker Box	Liu Wong	USA	10262	RATTC	2013-07-22	USA
TRAIH	Trails Head Gourmet Provisioners	Helvetius Nagy	USA	10262	RATTC	2013-07-22	USA
WHITC	White Clover Markets	Karl Jablonski	USA	10262	RATTC	2013-07-22	USA
GREAL	Great Lakes Food Market	Howard Snyder	USA	10269	WHITC	2013-07-31	USA
HUNGC	Hungry Coyote Import Store	Yoshi Latimer	USA	10269	WHITC	2013-07-31	USA
LAZYK	Lazy K Kountry Store	John Steel	USA	10269	WHITC	2013-07-31	USA
LETSS	Lets Stop N Shop	Jaime Yorres	USA	10269	WHITC	2013-07-31	USA
LONEP	Lonesome Pine Restaurant	Fran Wilson	USA	10269	WHITC	2013-07-31	USA
OLDWO	Old World Delicatessen	Rene Phillips	USA	10269	WHITC	2013-07-31	USA
RATTC	Rattlesnake Canyon Grocery	Paula Wilson	USA	10269	WHITC	2013-07-31	USA
SAVEA	Save-a-lot Markets	Jose Pavarotti	USA	10269	WHITC	2013-07-31	USA
SPLIR	Split Rail Beer & Ale	Art Braunschweiger	USA	10269	WHITC	2013-07-31	USA

Cartesian Product

- SQL must go through the Cartesian Product (which is an INTERIM answer set) row-by-row, and select only those rows where the EmployeeID from nwEmployees is equal to the EmployeeID from nwOrders
- Therefore, we must include the WHERE clause that describes this condition
- Failure to fully qualify a JOIN operation with a WHERE clause that matches all necessary keys will cause your answer set to include part or all of the Cartesian Product (which is mostly meaningless)
- The JOIN requires SQL to do a lot of work which consumes a lot of disk I/O and memory (= expensive)

Alternative join syntax

- SQL allows another syntax option for doing the JOIN.
- These queries are equivalent:

```
Select LastName, Firstname, count(OrderID) as 'Orders'  
      from nwEmployees E, nwOrders O  
     where E.EmployeeID = O.EmployeeID  
 GROUP BY LastName, FirstName  
 Order By 1
```

```
Select LastName, Firstname, count(OrderID) as 'Orders'  
      from nwEmployees E JOIN nwOrders O  
     on E.EmployeeID = O.EmployeeID  
 GROUP BY LastName, FirstName  
 Order By 1
```

Joining three or more tables

- Every PAIR of tables being joined must have a common key
- Every PAIR of common keys must have a condition stated in a WHERE clause or in the “ON” clause of the JOIN
- Otherwise, your JOIN is not fully qualified and will result in a Cartesian Product (meaningless output)

SQL Joins

Examples – Joining three tables

Create a report showing each employee and the total value of their orders sorted from highest value to lowest. (Order Value = UnitPrice * Quantity for each item on the order.)

```
Select LastName, FirstName,  
    sum(UnitPrice * Quantity) as 'OrderValue'  
from nwEmployees E  
JOIN nwOrders O on E.EmployeeID = O.EmployeeID  
JOIN nwOrderDetails D on O.OrderID = D.OrderID  
GROUP BY LastName, FirstName  
Order By 3 desc
```

Examples – Joining three tables

Same Query, Different Syntax

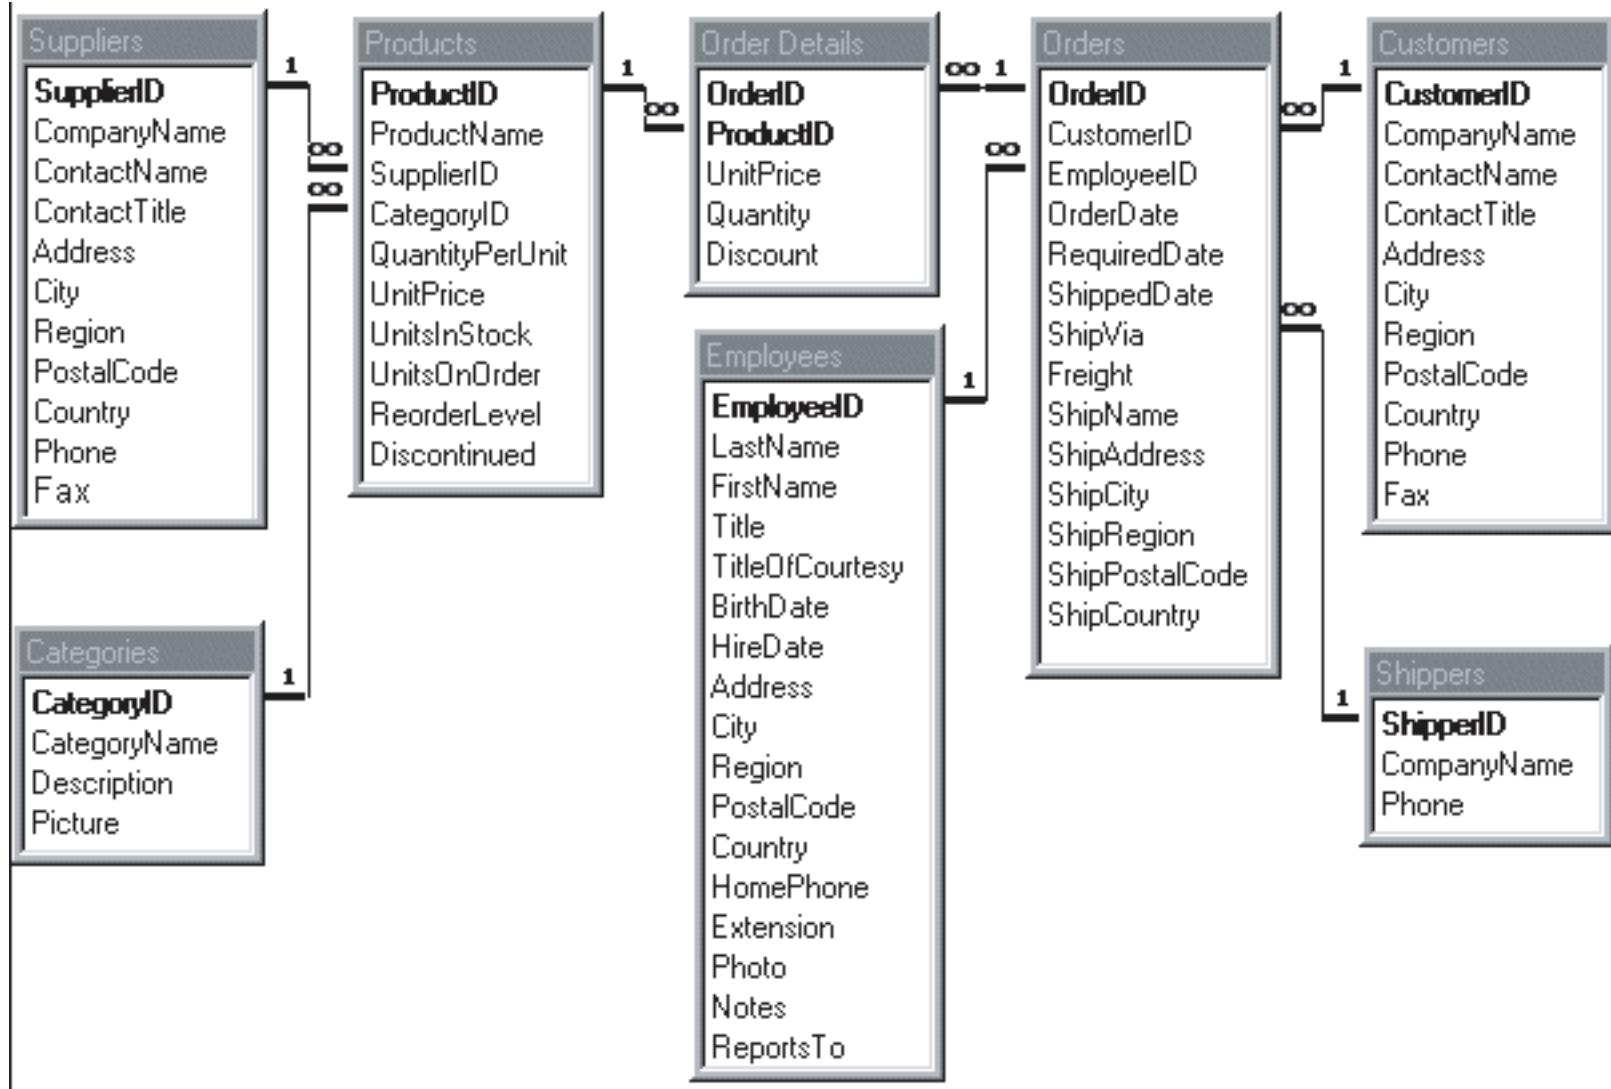
```
Select LastName, FirstName,  
    sum(UnitPrice * Quantity) as 'OrderValue'  
from nwEmployees E, nwOrders O,  
nwOrderDetails D  
where E.EmployeeID = O.EmployeeID  
    and O.OrderID = D.OrderID  
GROUP BY LastName, FirstName  
Order By 3 desc
```

Cartesian Product Error Example

- Consider a three-way join seeking data from 3 tables:
- Employees, Orders, and OrderDetails

“Provide a list of all employees and the total dollar value of each employee’s orders.”

SQL Joins



Cartesian Product Error Example

Provide a list of employees and the value of their orders.

Need to join Employees, Orders, and OrderDetails

Cartesian Product Error Example

```
Select LastName, FirstName,  
    sum(UnitPrice * Quantity) as 'OrderValue'  
from nwEmployees E, nwOrders O,  
        nwOrderDetails D  
where E.EmployeeID = O.EmployeeID  
GROUP BY LastName, FirstName  
Order By 3 desc
```

- **Join types**

Explicit inner join

```
SELECT *
FROM employee e
    INNER JOIN department d
        ON e.DepartmentID = d.DepartmentID
```

Implicit inner join:

```
SELECT * FROM employee e, department d
WHERE e.DepartmentID = d.DepartmentID
```

Left Outer Join

```
SELECT * FROM employees e  
LEFT OUTER JOIN department d  
ON e.DepartmentID = d.DepartmentID
```

Returns ALL rows from LEFT table and only matching rows from RIGHT table.

Right Outer Join

```
SELECT * FROM employees e  
RIGHT OUTER JOIN department d  
ON e.DepartmentID = d.DepartmentID
```

Returns ALL rows from RIGHT table and only matching rows from LEFT table.

Analysis using Outer Joins

1. Are there some customers who have no orders?
2. Are there orders in the nwOrders table that have an invalid reference to a Northwinds customer?
3. Are there orders in the nwOrders table that have an invalid reference to a Northwinds employee?

Analyzing Orders and Customers.

```
SELECT COUNT(customerid) FROM nwcustomers
```

- there are 87 customers in nwcustomers

```
SELECT COUNT(distinct customerid) FROM nworders
```

- there are 89 distinct customers in nworders

Is my data corrupt?

What's going on here...?

- **Invalid CustomerID?**

We want to find the orders in nworders whose customerID is NOT in nwcustomers

Method One: use a subquery

```
SELECT DISTINCT customerID  
    FROM nworders  
   WHERE customerID NOT IN (  
        SELECT customerID FROM nwcustomers) ;
```

Method Two: use an outer join

```
SELECT DISTINCT O.customerID  
    FROM nworders O LEFT OUTER JOIN nwcustomers C  
      ON O.customerID = C.customerID  
   WHERE C.customerID IS NULL
```

This shows us FOUR customers who have orders in nwOrders that have no matching row in nwCustomers

- **So what is the impact of this?**

(That is, what is the impact of having several orders with bad customerID's?)

But FIRST !

A note on UNION

- Assemble multiple queries connected by a UNION
- Combine multiple answer sets
- Each answer set must have the same number of columns
- Each column in all answer sets must have the same domain

EquiJoin:

```
SELECT C.customerID, CompanyName, COUNT(orderID)
  FROM nworders O JOIN nwcustomers C
    ON O.customerID = C.customerID
   GROUP BY C.customerID, CompanyName
UNION
SELECT 'Total', 'Total', COUNT(orderID)
  FROM nworders O JOIN nwcustomers C
    ON O.customerID = C.customerID;
```

This shows us a grand total of **785** orders

Outer Join:

```
SELECT DISTINCT O.customerID, CompanyName, COUNT(orderID)
  FROM nworders O LEFT OUTER JOIN nwcustomers C
    ON O.customerID = C.customerID
 GROUP BY O.customerID, CompanyName
UNION
SELECT 'Total', 'Total', COUNT(orderID)
  FROM nworders O LEFT OUTER JOIN nwcustomers C
    ON O.customerID = C.customerID;
```

This shows us a grand total of **829** orders

- So what is the impact of this (that several orders have bad customerID's)?
- This shows a discrepancy in the number of orders
 - 785 orders versus 829 orders

We could do a similar analysis comparing total orders dollar amounts (joining nwOrders, nwOrderDetails and nwCustomers)

Discrepancies in dollar amounts would fail an audit !!!

- We want to find the customers in nwcustmers who have no orders in nworders

Method One: use a subquery

```
SELECT customerID  
      FROM nwcustmers  
     WHERE customerID NOT IN (  
           SELECT DISTINCT customerID FROM nworders);
```

Method Two: use an outer join

```
SELECT DISTINCT C.customerID  
      FROM nworders O RIGHT OUTER JOIN nwcustmers C  
        ON O.customerID = C.customerID  
     WHERE O.customerID IS NULL
```

This shows us TWO customers who have no orders in nworders

- We want to find the employees in nwEmployees who have no orders in nwOrders

Method One: use a subquery

```
SELECT employeeID  
      FROM nwEmployees  
     WHERE employeeID NOT IN (  
           SELECT DISTINCT employeeID FROM nwOrders);
```

Method Two: use an outer join

```
SELECT DISTINCT e.employeeID  
      FROM nwEmployees E RIGHT OUTER JOIN nwOrders O  
        ON E.employeeID = O.employeeID  
     WHERE O.employeeID IS NULL
```

This shows us that all employees have some orders.

- We want to find any orders in nwOrders that have an invalid reference to nwEmployees

Method One: use a subquery

```
SELECT DISTINCT EmployeeID  
    FROM nworders  
   WHERE employeeID NOT IN (  
        SELECT employeeID FROM nwEmployees) );
```

Method Two: use an outer join

```
SELECT DISTINCT O.EmployeeID  
    FROM nwOrders O LEFT OUTER JOIN nwEmployees E  
      ON O.EmployeeID = E.employeeID  
 WHERE C.customerID IS NULL
```

This shows us that all orders have valid employeeIDs.

Join Execution Plans

- With any join, the database engine optimizer must calculate the most efficient query execution plan. There are three basic methods:

Nested Loop Join – when one join input table has a small number of rows and the other input table is large and indexed on the join key

Merge Join – when tables being joined are both sorted on the join key

Hash Join – when large, unsorted, non-indexed inputs are joined with an inner join with an "=" condition.

Join Execution Plans

- The “explain plan” option will show you an analysis of the execution plan calculation by the query optimizer.